



## MODULE 6: PREDICTIVE MODELING FOR TEMPORAL DATA

# CASE STUDY ACTIVITY TUTORIAL

### 6.1 New York City Case Study

# NewYorkCity\_taxi\_case\_study

November 26, 2017

## 1 New York City Taxi Ride Duration Prediction

In this case study, we will build a predictive model to predict the duration of taxi ride. We will do the following steps: \* Install the dependencies \* Load the data as pandas dataframe \* Define the outcome variable - the variable we are trying to predict. \* Build features with Deep Feature Synthesis using the [featuretools](#) package. We will start with simple features and incrementally improve the feature definitions and examine the accuracy of the system.

Allocate at least 2-3 hours to go through this case study end-to-end

## 2 Install Dependencies

If you have not done so already, download this repository from git. Once you have downloaded this archive, unzip it and cd into the directory from the command line. Next run the command `./install_osx.sh` if you are on a mac or `./install_linux.sh` if you are on linux. This should install all of the dependencies.

If you are on a windows machine, open the requirements.txt folder and make sure to install each of the dependencies listed (featuretools, jupyter, pandas, sklearn, numpy)

Once you have installed all of the dependencies, open this notebook. On Mac and Linux, navigate to the directory that you downloaded from git and run `jupyter notebook` to be taken to this notebook in your default web browser. When you open the `NewYorkCity_taxi_case_study.ipynb` file in the web browser, you can step through the code by clicking the Run button at the top of the page. If you have any questions for how to use Jupyter, refer to google or the discussion forum.

## 3 Running the Code

```
In [1]: import featuretools as ft
import utils
from utils import load_nyc_taxi_data, compute_features, preview, feature_importances
from sklearn.ensemble import GradientBoostingRegressor
from featuretools.primitives import (Weekend, Minute, Hour, Day, Week, Month,
                                     Weekday, Weekend, Count, Sum, Mean, Median,
                                     Std, Min, Max)

import numpy as np
ft.__version__
%load_ext autoreload
%autoreload 2
```

## 4 Step 1: Download and load the raw data as pandas dataframes

If you have not yet downloaded the data it can be downloaded from S3. Once you have downloaded the archive, unzip it and place the nyc-taxi-data folder in the same directory as this script.

```
In [2]: trips, pickup_neighborhoods, dropoff_neighborhoods = load_nyc_taxi_data()
        preview(trips, 10)
```

```
Out[2]:
```

	id	vendor_id		pickup_datetime		dropoff_datetime	\
0		0	2	2016-01-01 00:00:19	2016-01-01	00:06:31	
672146	672146	1	2016-04-29	07:01:31	2016-04-29	07:15:46	
672147	672147	2	2016-04-29	07:01:43	2016-04-29	07:09:15	
672148	672148	1	2016-04-29	07:01:46	2016-04-29	07:07:54	
672149	672149	2	2016-04-29	07:01:46	2016-04-29	07:06:48	
672150	672150	1	2016-04-29	07:01:59	2016-04-29	07:07:33	
672151	672151	2	2016-04-29	07:02:11	2016-04-29	07:15:24	
672152	672152	1	2016-04-29	07:02:11	2016-04-29	07:06:44	
672153	672153	2	2016-04-29	07:02:13	2016-04-29	07:08:36	
672154	672154	1	2016-04-29	07:02:16	2016-04-29	07:04:07	

  

	passenger_count	trip_distance	pickup_longitude	pickup_latitude	\
0	3	1.32	-73.961258	40.796200	
672146	1	3.30	-73.949951	40.784653	
672147	1	1.14	-73.967331	40.757370	
672148	1	1.10	-74.003082	40.727509	
672149	2	1.40	-73.990158	40.772350	
672150	1	1.20	-73.983681	40.746677	
672151	2	2.13	-73.994209	40.750999	
672152	1	1.00	-73.983276	40.770985	
672153	1	1.17	-73.980141	40.743168	
672154	1	0.50	-73.965973	40.765381	

  

	dropoff_longitude	dropoff_latitude	payment_type	trip_duration	\
0	-73.950050	40.787312	2	372.0	
672146	-73.982536	40.755470	1	855.0	
672147	-73.954277	40.765282	1	452.0	
672148	-73.984703	40.724377	1	368.0	
672149	-73.982147	40.759800	1	302.0	
672150	-73.971703	40.762463	2	334.0	
672151	-73.969391	40.761539	1	793.0	
672152	-73.980110	40.760666	1	273.0	
672153	-73.983391	40.754665	1	383.0	
672154	-73.970558	40.758724	1	111.0	

  

	pickup_neighborhood	dropoff_neighborhood
0	AH	C
672146	C	AA
672147	N	K

672148	AB	AC
672149	AR	AA
672150	AO	A
672151	D	AK
672152	AR	A
672153	Y	AA
672154	AK	N

The trips table has the following fields \* `id` which uniquely identifies the trip \* `vendor_id` is the taxi cab company - in our case study we have data from three different cab companies \* `pickup_datetime` the time stamp for pickup \* `dropoff_datetime` the time stamp for drop-off \* `passenger_count` the number of passengers for the trip \* `trip_distance` total distance of the trip in miles \* `pickup_longitude` the longitude for pickup \* `pickup_latitude` the latitude for pickup \* `dropoff_longitude` the longitude of dropoff \* `dropoff_latitude` the latitude of dropoff \* `payment_type` a numeric code signifying how the passenger paid for the trip. 1= Credit card 2= Cash 3= No charge 4= Dispute 5= Unknown 6= Voided \* `trip_duration` this is the duration we would like to predict using other fields \* `pickup_neighborhood` a one or two letter id of the neighborhood where the trip started \* `dropoff_neighborhood` a one or two letter id of the neighborhood where the trip ended

## 5 Step 2: Prepare the Data

Lets create entities and relationships. The three entities in this data are \* `trips` \* `pickup_neighborhoods` \* `dropoff_neighborhoods`

This data has the following relationships \* `pickup_neighborhoods` --> `trips` (one neighborhood can have multiple trips that start in it. This means `pickup_neighborhoods` is the parent\_entity and `trips` is the child entity) \* `dropoff_neighborhoods` --> `trips` (one neighborhood can have multiple trips that end in it. This means `dropoff_neighborhoods` is the parent\_entity and `trips` is the child entity)

In , we specify the list of entities and relationships as follows:

```
In [3]: entities = {
    "trips": (trips, "id", 'pickup_datetime' ),
    "pickup_neighborhoods": (pickup_neighborhoods, "neighborhood_id"),
    "dropoff_neighborhoods": (dropoff_neighborhoods, "neighborhood_id"),
}

relationships = [("pickup_neighborhoods", "neighborhood_id",
    "trips", "pickup_neighborhood"),
    ("dropoff_neighborhoods", "neighborhood_id",
    "trips", "dropoff_neighborhood")]
```

Next, we specify the cutoff time for each instance of the target\_entity, in this case `trips`. This timestamp represents the last time data can be used for calculating features by DFS. In this scenario, that would be the pickup time because we would like to make the duration prediction using data before the trip starts.

For the purposes of the case study, we choose to only select trips that started after January 12th, 2016.

```
In [4]: cutoff_time = trips[['id', 'pickup_datetime']]
        cutoff_time = cutoff_time[cutoff_time['pickup_datetime'] > "2016-01-12"]
        preview(cutoff_time, 10)
```

```
Out[4]:
```

	id	pickup_datetime
56311	56311	2016-01-12 00:00:25
698765	698765	2016-05-03 18:54:53
698766	698766	2016-05-03 18:55:37
698767	698767	2016-05-03 18:55:38
698768	698768	2016-05-03 18:55:49
698769	698769	2016-05-03 18:55:58
698770	698770	2016-05-03 18:56:22
698771	698771	2016-05-03 18:56:24
698772	698772	2016-05-03 18:56:51
698773	698773	2016-05-03 18:56:56

## 6 Step 3: Create baseline features using Deep Feature Synthesis

Instead of manually creating features, such as "month of pickup datetime", we can let DFS come up with them automatically. It does this by \* interpreting the variable types of the columns e.g categorical, numeric and others \* matching the columns to the primitives that can be applied to their variable types \* creating features based on these matches

## 7 Create transform features using transform primitives

As we described in the video, features fall into two major categories, transform and aggregate. In featuretools, we can create transform features by specifying transform primitives. Below we specify a transform primitive called weekend and here is what it does:

- It can be applied to any datetime column in the data.
- For each entry in the column, it assess if it is a weekend and returns a boolean.

In this specific data, there are two datetime columns pickup\_datetime and dropoff\_datetime. The tool automatically creates features using the primitive and these two columns as shown below.

```
In [5]: trans_primitives = [Weekend]

features = ft.dfs(entities=entities,
                  relationships=relationships,
                  target_entity="trips",
                  trans_primitives=trans_primitives,
                  agg_primitives=[],
                  ignore_variables={"trips": ["pickup_latitude",
                                              "pickup_longitude",
                                              "dropoff_latitude",
                                              "dropoff_longitude"]},
                  features_only=True)
```

If you're interested about parameters to DFS such as *ignore\_variables*, you can learn more about these parameters [here](#)

Here are the features created.

```
In [6]: print "Number of features: %d" % len(features)
        features
```

Number of features: 13

```
Out[6]: [<Feature: vendor_id>,
        <Feature: passenger_count>,
        <Feature: dropoff_neighborhood>,
        <Feature: payment_type>,
        <Feature: pickup_neighborhood>,
        <Feature: trip_duration>,
        <Feature: trip_distance>,
        <Feature: dropoff_neighborhoods.longitude>,
        <Feature: dropoff_neighborhoods.latitude>,
        <Feature: pickup_neighborhoods.longitude>,
        <Feature: pickup_neighborhoods.latitude>,
        <Feature: IS_WEEKEND(pickup_datetime)>,
        <Feature: IS_WEEKEND(dropoff_datetime)>]
```

Now let's compute the features.

```
In [7]: feature_matrix = compute_features(features, cutoff_time)
```

Progress: 100%|| 5/5 [00:16<00:00, 3.20s/cutoff time]

Finishing computing...

```
In [8]: preview(feature_matrix, 5)
```

```
Out[8]:
```

	IS_WEEKEND(dropoff_datetime)	trip_distance	passenger_count	\
id				
56311	False	1.61	1	
691284	False	0.61	2	
691285	False	0.88	2	
691286	False	1.90	1	
691288	False	1.00	1	

  

	dropoff_neighborhood = AD	dropoff_neighborhood = A	\
id			
56311	0	0	
691284	0	0	
691285	0	0	
691286	0	0	
691288	0	0	

	dropoff_neighborhood = AA	dropoff_neighborhood = D \
id		
56311	0	0
691284	0	0
691285	0	0
691286	1	0
691288	0	0

	dropoff_neighborhood = AR	dropoff_neighborhood = C \
id		
56311	0	0
691284	0	0
691285	0	0
691286	0	0
691288	0	0

	dropoff_neighborhood = 0	...	\
id		...	
56311	0	...	
691284	0	...	
691285	0	...	
691286	0	...	
691288	0	...	

	pickup_neighborhood = AD	pickup_neighborhood = AA \
id		
56311	0	0
691284	0	0
691285	0	0
691286	0	0
691288	0	0

	pickup_neighborhood = D	pickup_neighborhood = A \
id		
56311	0	0
691284	0	0
691285	0	0
691286	0	0
691288	0	0

	pickup_neighborhood = AR	pickup_neighborhood = AK \
id		
56311	0	0
691284	0	0
691285	0	0
691286	0	0
691288	0	0



```

pickup_neighborhood = A0 pickup_neighborhood = N \
id
56311                0                0
691284                0                0
691285                0                0
691286                0                0
691288                1                0

pickup_neighborhood = R pickup_neighborhood = 0
id
56311                0                0
691284                0                0
691285                0                0
691286                0                0
691288                0                0

[5 rows x 31 columns]

```

## 8 Step 4: Build the Model

To build a model, we \* Seperate the data into a porition for training (75% in this case) and a portion for testing \* Get the log of the trip duration so that a more linear relationship can be found. \* Train a model using a GradientBoostingRegressor

```

In [9]: # separates the whole feature matrix into train data feature matrix,
        # train data labels, and test data feature matrix
X_train, y_train, X_test, y_test = utils.get_train_test_fm(feature_matrix,.75)
y_train = np.log(y_train+1)
y_test = np.log(y_test+1)

In [10]: model = GradientBoostingRegressor(verbose=True)
         model.fit(X_train, y_train)
         model.score(X_test, y_test)

```

Iter	Train Loss	Remaining Time
1	0.4925	2.50m
2	0.4333	2.47m
3	0.3843	2.43m
4	0.3446	2.47m
5	0.3119	2.44m
6	0.2852	2.40m
7	0.2634	2.37m
8	0.2454	2.33m
9	0.2305	2.30m
10	0.2183	2.42m
20	0.1666	2.32m
30	0.1558	1.99m



40	0.1514	1.70m
50	0.1488	1.34m
60	0.1472	1.05m
70	0.1458	46.77s
80	0.1448	30.91s
90	0.1440	15.44s
100	0.1433	0.00s

Out[10]: 0.72200175704571445

## 9 Step 5: Adding more Transform Primitives

- Add Minute, Hour, Week, Month, Weekday , etc primitives
- All these transform primitives apply to datetime columns

```
In [11]: trans_primitives = [Minute, Hour, Day, Week, Month, Weekday, Weekend]
```

```
features = ft.dfs(entities=entities,
                  relationships=relationships,
                  target_entity="trips",
                  trans_primitives=trans_primitives,
                  agg_primitives=[],
                  ignore_variables={"trips": ["pickup_latitude",
                                             "pickup_longitude",
                                             "dropoff_latitude",
                                             "dropoff_longitude"]},
                  features_only=True)
```

```
In [12]: print "Number of features: %d" % len(features)
features
```

Number of features: 25

```
Out[12]: [<Feature: passenger_count>,
<Feature: dropoff_neighborhood>,
<Feature: payment_type>,
<Feature: vendor_id>,
<Feature: pickup_neighborhood>,
<Feature: trip_duration>,
<Feature: trip_distance>,
<Feature: DAY(pickup_datetime)>,
<Feature: dropoff_neighborhoods.latitude>,
<Feature: WEEK(dropoff_datetime)>,
<Feature: HOUR(pickup_datetime)>,
<Feature: WEEKDAY(dropoff_datetime)>,
<Feature: WEEKDAY(pickup_datetime)>]
```

```

<Feature: MONTH(pickup_datetime)>,
<Feature: WEEK(pickup_datetime)>,
<Feature: DAY(dropoff_datetime)>,
<Feature: MONTH(dropoff_datetime)>,
<Feature: pickup_neighborhoods.latitude>,
<Feature: HOUR(dropoff_datetime)>,
<Feature: pickup_neighborhoods.longitude>,
<Feature: IS_WEEKEND(pickup_datetime)>,
<Feature: MINUTE(pickup_datetime)>,
<Feature: MINUTE(dropoff_datetime)>,
<Feature: dropoff_neighborhoods.longitude>,
<Feature: IS_WEEKEND(dropoff_datetime)>]

```

Now let's compute the features.

```
In [13]: feature_matrix = compute_features(features, cutoff_time)
```

```
Progress: 100%|| 5/5 [00:21<00:00, 4.24s/cutoff time]
Finishing computing...
```

```
In [14]: preview(feature_matrix, 10)
```

```
Out[14]:
```

	WEEKDAY(dropoff_datetime)	dropoff_neighborhoods.latitude \
id		
56311	1	40.721435
691284	0	40.721435
691285	0	40.785005
691286	0	40.757707
691288	0	40.761087
691289	0	40.761492
691290	0	40.764723
691291	0	40.776270
691292	0	40.764723
691293	0	40.766488

  

	MINUTE(dropoff_datetime)	WEEK(dropoff_datetime)	passenger_count \
id			
56311	11	2	1
691284	24	18	2
691285	27	18	2
691286	48	18	1
691288	30	18	1
691289	55	18	1
691290	26	18	1
691291	37	18	1
691292	39	18	1
691293	34	18	2

	trip_duration	hour(pickup_datetime)	pickup_neighborhoods.latitude	\
id				
56311	645.0	0	40.720245	
691284	160.0	12	40.729652	
691285	295.0	12	40.776270	
691286	1573.0	12	40.742531	
691288	404.0	12	40.747126	
691289	1906.0	12	40.721435	
691290	156.0	12	40.764723	
691291	827.0	12	40.766809	
691292	883.0	12	40.752186	
691293	592.0	12	40.775299	

	vendor_id	dropoff_neighborhoods.longitude	...	\
id			...	
56311	2	-73.998366	...	
691284	2	-73.998366	...	
691285	2	-73.976050	...	
691286	1	-73.986446	...	
691288	1	-73.995736	...	
691289	2	-73.975899	...	
691290	1	-73.966696	...	
691291	1	-73.982322	...	
691292	1	-73.966696	...	
691293	2	-73.983998	...	

	dropoff_neighborhood = D	dropoff_neighborhood = AR	\
id			
56311	0	0	
691284	0	0	
691285	0	0	
691286	0	0	
691288	0	0	
691289	0	0	
691290	0	0	
691291	0	0	
691292	0	0	
691293	0	1	

	dropoff_neighborhood = C	dropoff_neighborhood = 0	\
id			
56311	0	0	
691284	0	0	
691285	0	0	
691286	0	0	
691288	0	0	
691289	0	0	
691290	0	0	

691291	0	0
691292	0	0
691293	0	0

  

	dropoff_neighborhood = N	dropoff_neighborhood = A0 \
id		
56311	0	0
691284	0	0
691285	0	0
691286	0	0
691288	0	0
691289	0	0
691290	0	0
691291	0	0
691292	0	0
691293	0	0

  

	dropoff_neighborhood = AK	HOUR(dropoff_datetime) \
id		
56311	0	0
691284	0	12
691285	0	12
691286	0	12
691288	0	12
691289	0	12
691290	1	12
691291	0	12
691292	1	12
691293	0	12

  

	IS_WEEKEND(dropoff_datetime)	trip_distance
id		
56311	False	1.61
691284	False	0.61
691285	False	0.88
691286	False	1.90
691288	False	1.00
691289	False	3.24
691290	False	0.10
691291	False	1.60
691292	False	1.50
691293	False	1.89

[10 rows x 43 columns]

## 10 Step 6: Build the new model

```
In [15]: # separates the whole feature matrix into train data feature matrix,
# train data labels, and test data feature matrix
X_train, y_train, X_test, y_test = utils.get_train_test_fm(feature_matrix,.75)
y_train = np.log(y_train+1)
y_test = np.log(y_test+1)

In [16]: model = GradientBoostingRegressor(verbose=True)
model.fit(X_train,y_train)
model.score(X_test,y_test)
```

Iter	Train Loss	Remaining Time
1	0.4925	3.86m
2	0.4333	3.84m
3	0.3843	3.77m
4	0.3444	3.73m
5	0.3117	3.78m
6	0.2848	3.75m
7	0.2620	3.70m
8	0.2435	3.66m
9	0.2282	3.60m
10	0.2152	3.56m
20	0.1588	3.11m
30	0.1415	2.54m
40	0.1332	2.09m
50	0.1283	1.67m
60	0.1252	1.30m
70	0.1227	59.48s
80	0.1207	40.02s
90	0.1191	19.87s
100	0.1177	0.00s

```
Out[16]: 0.7755608981558122
```

## 11 Step 7: Add Aggregation Primitives

Now let's add aggregation primitives. These primitives will generate features for the parent entities pickup\_neighborhoods, and dropoff\_neighborhood and then add them to the trips entity, which is the entity for which we are trying to make prediction.

```
In [17]: trans_primitives = [Minute, Hour, Day, Week, Month, Weekday, Weekend]
aggregation_primitives = [Count, Sum, Mean, Median, Std, Max, Min]

features = ft.dfs(entities=entities,
                  relationships=relationships,
                  target_entity="trips",
```

```

trans_primitives=trans_primitives,
agg_primitives=aggregation_primitives,
ignore_variables={"trips": ["pickup_latitude",
                             "pickup_longitude",
                             "dropoff_latitude",
                             "dropoff_longitude"]},

features_only=True)

```

```

In [18]: print "Number of features: %d" % len(features)
         features

```

Number of features: 63

```

Out[18]: [<Feature: passenger_count>,
<Feature: dropoff_neighborhood>,
<Feature: payment_type>,
<Feature: vendor_id>,
<Feature: pickup_neighborhood>,
<Feature: trip_duration>,
<Feature: trip_distance>,
<Feature: DAY(pickup_datetime)>,
<Feature: dropoff_neighborhoods.latitude>,
<Feature: WEEK(dropoff_datetime)>,
<Feature: HOUR(pickup_datetime)>,
<Feature: WEEKDAY(dropoff_datetime)>,
<Feature: WEEKDAY(pickup_datetime)>,
<Feature: MONTH(pickup_datetime)>,
<Feature: WEEK(pickup_datetime)>,
<Feature: pickup_neighborhoods.latitude>,
<Feature: DAY(dropoff_datetime)>,
<Feature: MONTH(dropoff_datetime)>,
<Feature: HOUR(dropoff_datetime)>,
<Feature: pickup_neighborhoods.longitude>,
<Feature: dropoff_neighborhoods.longitude>,
<Feature: IS_WEEKEND(pickup_datetime)>,
<Feature: MINUTE(pickup_datetime)>,
<Feature: MINUTE(dropoff_datetime)>,
<Feature: IS_WEEKEND(dropoff_datetime)>,
<Feature: dropoff_neighborhoods.SUM(trips.trip_duration)>,
<Feature: pickup_neighborhoods.STD(trips.trip_duration)>,
<Feature: pickup_neighborhoods.MEDIAN(trips.trip_duration)>,
<Feature: dropoff_neighborhoods.STD(trips.trip_distance)>,
<Feature: pickup_neighborhoods.MEAN(trips.trip_duration)>,
<Feature: pickup_neighborhoods.MIN(trips.trip_duration)>,
<Feature: dropoff_neighborhoods.MEDIAN(trips.trip_duration)>,
<Feature: pickup_neighborhoods.MEAN(trips.trip_distance)>,
<Feature: pickup_neighborhoods.SUM(trips.trip_duration)>,

```

```

<Feature: pickup_neighborhoods.MIN(trips.trip_distance)>,
<Feature: pickup_neighborhoods.MAX(trips.trip_distance)>,
<Feature: dropoff_neighborhoods.MEAN(trips.passenger_count)>,
<Feature: dropoff_neighborhoods.MEDIAN(trips.passenger_count)>,
<Feature: dropoff_neighborhoods.MAX(trips.trip_duration)>,
<Feature: pickup_neighborhoods.MAX(trips.trip_duration)>,
<Feature: dropoff_neighborhoods.STD(trips.trip_duration)>,
<Feature: dropoff_neighborhoods.STD(trips.passenger_count)>,
<Feature: dropoff_neighborhoods.MIN(trips.passenger_count)>,
<Feature: dropoff_neighborhoods.MEAN(trips.trip_duration)>,
<Feature: pickup_neighborhoods.STD(trips.passenger_count)>,
<Feature: pickup_neighborhoods.SUM(trips.trip_distance)>,
<Feature: pickup_neighborhoods.COUNT(trips)>,
<Feature: pickup_neighborhoods.STD(trips.trip_distance)>,
<Feature: dropoff_neighborhoods.SUM(trips.trip_distance)>,
<Feature: pickup_neighborhoods.MAX(trips.passenger_count)>,
<Feature: pickup_neighborhoods.MIN(trips.passenger_count)>,
<Feature: dropoff_neighborhoods.MEDIAN(trips.trip_distance)>,
<Feature: pickup_neighborhoods.MEDIAN(trips.trip_distance)>,
<Feature: dropoff_neighborhoods.COUNT(trips)>,
<Feature: pickup_neighborhoods.MEDIAN(trips.passenger_count)>,
<Feature: pickup_neighborhoods.SUM(trips.passenger_count)>,
<Feature: dropoff_neighborhoods.MEAN(trips.trip_distance)>,
<Feature: dropoff_neighborhoods.SUM(trips.passenger_count)>,
<Feature: dropoff_neighborhoods.MIN(trips.trip_duration)>,
<Feature: dropoff_neighborhoods.MAX(trips.trip_distance)>,
<Feature: dropoff_neighborhoods.MIN(trips.trip_distance)>,
<Feature: pickup_neighborhoods.MEAN(trips.passenger_count)>,
<Feature: dropoff_neighborhoods.MAX(trips.passenger_count)>]

```

```
In [19]: feature_matrix = compute_features(features, cutoff_time)
```

```
Progress: 100%|| 5/5 [00:53<00:00, 10.23s/cutoff time]
```

```
Finishing computing...
```

```
In [20]: preview(feature_matrix, 10)
```

```
Out[20]:      dropoff_neighborhoods.MAX(trips.trip_duration) \
id
56311      3572.0
691284      3603.0
691285      3602.0
691286      3606.0
691288      3580.0
691289      3606.0
691290      3580.0
691291      3604.0
691292      3580.0
```



691293	3587.0
--------	--------

pickup_neighborhoods.MEDIAN(trips.passenger_count) \	
id	
56311	1.0
691284	1.0
691285	1.0
691286	1.0
691288	1.0
691289	1.0
691290	1.0
691291	1.0
691292	1.0
691293	1.0

pickup_neighborhoods.MEDIAN(trips.trip_distance) \	
id	
56311	2.40
691284	1.60
691285	1.60
691286	1.49
691288	1.40
691289	1.90
691290	1.30
691291	1.63
691292	1.49
691293	1.37

HOUR(dropoff_datetime) dropoff_neighborhoods.COUNT(trips) \		
id		
56311	0	1396.0
691284	12	16736.0
691285	12	19017.0
691286	12	28805.0
691288	12	16985.0
691289	12	31541.0
691290	12	21894.0
691291	12	21272.0
691292	12	21894.0
691293	12	24592.0

DAY(pickup_datetime) pickup_neighborhoods.latitude \		
id		
56311	12	40.720245
691284	2	40.729652
691285	2	40.776270
691286	2	40.742531
691288	2	40.747126

691289	2	40.721435
691290	2	40.764723
691291	2	40.766809
691292	2	40.752186
691293	2	40.775299

```

pickup_neighborhoods.SUM(trips.passenger_count) \
id
56311                2283.0
691284                34521.0
691285                36299.0
691286                31158.0
691288                43543.0
691289                30913.0
691290                43212.0
691291                32656.0
691292                57862.0
691293                39612.0

```

```

pickup_neighborhoods.STD(trips.trip_distance) \
id
56311                2.517060
691284                2.099009
691285                2.111243
691286                2.137177
691288                2.382449
691289                4.278882
691290                1.846378
691291                2.206183
691292                2.488034
691293                1.904818

```

```

dropoff_neighborhoods.SUM(trips.passenger_count) \
id
56311                2375.0
691284                28154.0
691285                31836.0
691286                49208.0
691288                28197.0
691289                52591.0
691290                36175.0
691291                35282.0
691292                36175.0
691293                41249.0

```

```

... \
id ...
56311 ...

```

691284	...
691285	...
691286	...
691288	...
691289	...
691290	...
691291	...
691292	...
691293	...

	pickup_neighborhoods.STD(trips.passenger_count) \
id	
56311	1.331649
691284	1.310235
691285	1.315396
691286	1.330198
691288	1.319326
691289	1.315238
691290	1.315462
691291	1.332742
691292	1.306561
691293	1.341478

	MONTH(pickup_datetime) \
id	
56311	1
691284	5
691285	5
691286	5
691288	5
691289	5
691290	5
691291	5
691292	5
691293	5

	pickup_neighborhoods.MEAN(trips.trip_duration)	WEEK(pickup_datetime) \
id		
56311	740.870871	2
691284	753.813680	18
691285	681.405688	18
691286	682.624440	18
691288	714.648716	18
691289	818.141251	18
691290	637.726834	18
691291	707.024093	18
691292	749.696305	18
691293	670.677993	18

	dropoff_neighborhoods.MEAN(trips.trip_distance) \
id	
56311	2.495358
691284	2.338798
691285	2.176976
691286	2.365290
691288	2.067381
691289	2.102551
691290	1.732215
691291	2.061938
691292	1.732215
691293	2.200316

	MONTH(dropoff_datetime)	payment_type	MINUTE(dropoff_datetime) \
id			
56311	1	1	11
691284	5	1	24
691285	5	1	27
691286	5	1	48
691288	5	1	30
691289	5	1	55
691290	5	2	26
691291	5	1	37
691292	5	1	39
691293	5	1	34

	WEEK(dropoff_datetime) \
id	
56311	2
691284	18
691285	18
691286	18
691288	18
691289	18
691290	18
691291	18
691292	18
691293	18

	dropoff_neighborhoods.MAX(trips.passenger_count)
id	
56311	6.0
691284	6.0
691285	6.0
691286	6.0
691288	6.0
691289	6.0

691290	6.0
691291	6.0
691292	6.0
691293	6.0

[10 rows x 81 columns]

## 12 Step 8: Build the new model

```
In [21]: # separates the whole feature matrix into train data feature matrix,
# train data labels, and test data feature matrix
X_train, y_train, X_test, y_test = utils.get_train_test_fm(feature_matrix,.75)
y_train = np.log(y_train+1)
y_test = np.log(y_test+1)
```

```
In [22]: # note: this may take up to 30 minutes to run
model = GradientBoostingRegressor(verbose=True)
model.fit(X_train, y_train)
```

Iter	Train Loss	Remaining Time
1	0.4925	11.57m
2	0.4333	10.96m
3	0.3843	10.33m
4	0.3444	10.06m
5	0.3117	9.87m
6	0.2848	9.61m
7	0.2620	9.71m
8	0.2435	9.58m
9	0.2282	9.47m
10	0.2152	9.30m
20	0.1585	7.99m
30	0.1420	6.74m
40	0.1332	5.57m
50	0.1271	4.47m
60	0.1238	3.51m
70	0.1211	2.63m
80	0.1191	1.72m
90	0.1176	50.97s
100	0.1163	0.00s

```
Out[22]: GradientBoostingRegressor(alpha=0.9, criterion='friedman_mse', init=None,
learning_rate=0.1, loss='ls', max_depth=3, max_features=None,
max_leaf_nodes=None, min_impurity_decrease=0.0,
min_impurity_split=None, min_samples_leaf=1,
min_samples_split=2, min_weight_fraction_leaf=0.0,
n_estimators=100, presort='auto', random_state=None,
subsample=1.0, verbose=True, warm_start=False)
```

## 13 Step 9: Evaluate on test data

```
In [23]: model.score(X_test,y_test)
```

```
Out[23]: 0.77808885638984859
```

we can also make predictions using our model

```
In [24]: y_pred = model.predict(X_test)
         y_pred = np.exp(y_pred) - 1 # undo the log we took earlier
         y_pred[5:]
```

```
Out[24]: array([ 557.67992664,  590.2602792 , 1497.39684679, ..., 1063.48382696,
                1800.89361932,  739.60249439])
```

## 14 Additional Analysis

Let's look at how important each feature was for the model.

```
In [25]: feature_importances(model, feature_matrix.columns, n=15)
```

```
1: Feature: trip_distance, 0.314
2: Feature: HOUR(pickup_datetime), 0.126
3: Feature: HOUR(dropoff_datetime), 0.089
4: Feature: WEEKDAY(pickup_datetime), 0.052
5: Feature: dropoff_neighborhoods.latitude, 0.046
6: Feature: dropoff_neighborhoods.longitude, 0.036
7: Feature: dropoff_neighborhoods.STD(trips.trip_distance), 0.027
8: Feature: dropoff_neighborhoods.MIN(trips.passenger_count), 0.022
9: Feature: dropoff_neighborhoods.MEDIAN(trips.trip_duration), 0.022
10: Feature: pickup_neighborhoods.MEDIAN(trips.trip_distance), 0.021
11: Feature: IS_WEEKEND(pickup_datetime), 0.021
12: Feature: WEEKDAY(dropoff_datetime), 0.020
13: Feature: WEEK(pickup_datetime), 0.019
14: Feature: dropoff_neighborhoods.MEAN(trips.trip_duration), 0.019
15: Feature: MONTH(dropoff_datetime), 0.018
```