



## MODULE 4: RECOMMENDATION SYSTEMS

# CASE STUDY ACTIVITY TUTORIAL

## CASE STUDY 3 – NEW PRODUCT RECOMMENDATION

# CASE STUDY ACTIVITY TUTORIAL

## CASE STUDY 3 – NEW PRODUCT RECOMMENDATION

**Problem:** Make new Product Recommendations (e.g. Amazon.com)

- **Dataset:** Amazon Reviews data (<http://jmcauley.ucsd.edu/data/amazon/>)
  - You will need to seek permission to use any/all of the datasets in this repository
  - The repository has several datasets. You can choose any. For this example we will use the Electronics dataset.
  - The host page has several pointers to scripts and other examples that can help with parsing the datasets.
  - The data set consists of:
    - 7,824,482 Ratings (1-5) for Electronics products.
    - Other metadata about products. Please see the description of the fields available on the webpage cited above.
  - For convenience of future use, parse the raw data file (using Python, for example) and extract the following fields: 'product/productId' as **prod\_id**, 'product/title' as **prod\_name**, 'review/userId' as **user id**, 'review/score' as **rating**
  - Save these to a tab separated file. We will refer to this file as **product\_ratings.csv**
- **Read/View the Dataset:**
  - The first task is to explore the dataset. You can do so using a programming environment of your choice, e.g. Python or R.
  - In R, you can read the data by simply calling the read.table() function:

```
data = read.csv('product_ratings.csv')
```

    - You can rename the column names as desired:

```
colnames(data) = c("prod_id", "prod_name", "user_id", "rating")
```
    - Now you can look at the data properties by using:

```
str(data)
```

```
summary(data)
```
    - Plot a histogram of the data:

```
hist(data$rating)
```
  - In Python, you can convert the data to a pandas dataframe to organize the dataset (<http://pandas.pydata.org/>)  
For plotting in Python, you can use Matplotlib: <http://matplotlib.org/>
  - The dataset sparsity can be calculated as:  
$$\text{Sparsity} = \frac{\text{Number of Ratings in the Dataset}}{(\text{Number of Products} \times \text{Number of Users})} \times 100\%$$

- **Sub-setting the data:**

- If you want the data to be less sparse, for example, a good way to achieve that is to subset the data where you only select Users/Products that have at least a certain number of observations in the dataset.
- In R, for example, if you wanted to subset the data such that only users with 50 or more ratings remained, you would do the following:

```
data = data[ data$user_id %in%
names(table(data$user_id))[table(data$user_id) > 50] , ]
```

- **Recommenders:**

- If you want to build your own Recommenders from scratch, you can consult the vast amounts of academic literature available freely. There are also several self-help guides which can be useful, such as these:

<http://www.salemmarafi.com/code/collaborative-filtering-r/> ,

<http://blogs.gartner.com/martin-kihn/how-to-build-a-recommender-system-in-python/>

- On the other hand, why build a recommender from scratch when there is a vast array of publicly available Recommenders (in all sorts of programming environments) ready for use? Some examples are:

- RecommenderLab in R (<https://cran.r-project.org/web/packages/recommenderlab/vignettes/recommenderlab.pdf>) ,
- Graphlab-Create for Python (has a free license for persona and academic use)(<https://dato.com/products/create/docs/graphlab.toolkits.recommender.html>),
- Apache Spark's Recommendation module (<https://spark.apache.org/docs/1.4.0/api/python/pyspark.mllib.html#module-pyspark.mllib.recommendation>),
- Apache Mahout (<https://mahout.apache.org/users/recommender/userbased-5-minutes.html>)

- **Splitting Data Randomly (Train/Test):**

- A random split can be created in R and Pandas (Python) .
- In R, you can do the following to create a 70/30 split for Train/Test:

```
library(caTools)
spl = sample.split(data$rating, 0.7)
train = subset(data, spl == TRUE)
test = subset(data, spl == FALSE)
```

- In Pandas, using the SciKit-Learn library:
 

```
import pandas as pd
import numpy as np
from sklearn.cross_validation import train_test_split

# assuming pdf is the pandas dataframe with the data
train, test = train_test_split(pdf, test_size = 0.3)
```
- Alternatively, one can use the Recommender libraries (discussed earlier) to create the data splits.
  - For RecommenderLab in R, the documentation in Section 5.6 provides examples that will allow random data splits (<https://cran.r-project.org/web/packages/recommenderlab/vignettes/recommenderlab.pdf>)
  - Graphlab's Sframe objects also have a random\_split function which works similarly ([https://dato.com/products/create/docs/generated/graphlab.SFrame.random\\_split.html](https://dato.com/products/create/docs/generated/graphlab.SFrame.random_split.html))
- **Popularity Recommender:**
  - The RecommenderLab in R, for example, provides a popularity recommender out of the box. Section 5.5 of the RecommenderLab guide provides examples and sample code to help do this: <https://cran.r-project.org/web/packages/recommenderlab/vignettes/recommenderlab.pdf>
  - GraphLab-Create also provides a Popularity Recommender (in python). If the dataset is in Pandas, it can easily integrate with GraphLab's Sframe datatype as noted here: <https://dato.com/products/create/docs/generated/graphlab.SFrame.html>. Some more information on the Popularity Recommender and its usage is provided here: [https://dato.com/products/create/docs/generated/graphlab.recommender.popularity\\_recommender.PopularityRecommender.html](https://dato.com/products/create/docs/generated/graphlab.recommender.popularity_recommender.PopularityRecommender.html)
- **Collaborative Filtering:**
  - Most recommender libraries will provide an implementation for Collaborative Filtering methods. The RecommenderLab in R and GraphLab in Python both provide implementations of Collaborative Filtering methods.
    - In RecommenderLab, use the "UBCF" (user-based collaborative filtering) to train the model.
    - In GraphLab, use the "Factorization Recommender" ([https://dato.com/products/create/docs/generated/graphlab.recommender.factorization\\_recommender.FactorizationRecommender.html](https://dato.com/products/create/docs/generated/graphlab.recommender.factorization_recommender.FactorizationRecommender.html))

- Often, a regularization parameter is used with these models. The best value for this regularization parameter is chosen using a Validation set. Here is how this can be done:
  - Split the Test set further 75%/25% in to Train/Validation sets.
  - Now we have three sets: Train, Validation, Test.
  - Train several models, each using a different value of the regularization parameter (usually in the range:  $(1e-5, 1e-1)$ ).
  - Use the Validation set to determine which model results in the lowest RMSE.
  - Use the regularization value that corresponds to the lowest Validation set RMSE.
  - Finally, with that parameter value fixed, use the trained model to get a final RMSE value on the Test set.
  - In R and Python, it can also help plotting the Validation set RMSE values vs the Regularization parameter values to determine the best one.
- **Evaluation (RMSE):** Once the model is trained on the Training data, it can be used to compute the error (RMSE) on predictions made on the Test data.
  - RecommenderLab in R uses the **predict()** and **calcPredictionAccuracy()** functions to compute the predictions (based on the trained model) and evaluate RMSE (and MSE and MAE).
  - Graphlab in Python also has a **predict()** function to get predictions. It provides a suite of functions to evaluate metrics such as rmse (**evaluation.rmse()**, for example).
- **Top-K Recommendations:**
  - Since our goal is to recommend new products to each user based on his/her habits, we will recommend K new products.
  - Based on scores assigned to User-Item pairs, each recommender algorithm makes available functions that will provide a sorted list of top-K items most highly recommended for each user (from among those items **not** already rated by the user).
    - In RecommenderLab, the parameter **type='topNlist'** to the **evaluate()** function will produce such a list.
    - In GraphLab, the **recommend(K)** function for each type of recommender will do the same.